



Programming language & software security

COSC312 / COSC412


Learning objectives

- **Choice of programming language** can affect security
 - ... although choice of PL is almost certainly not a panacea
- High level view of causes of **PL/software security issues**
 - Provide a roadmap into which to fit common attack types
- Many **software security checks** can be regimented
 - Also that there are many, many vectors for security attacks!

Typical computing machine model

- Programming language security depends on machine:
 - we'll assume typical **von Neumann architecture** depicted

- **CPU** runs imperative code
- **I/O** devices (input/output)
- **Memory** (code and data)
 - hierarchy of memory levels
- First we'll focus on CPU behaviour



Security from machine model's perspective

- Risks in terms of the CPU going awry:
 - **Space**—CPU interacts with memory in unintended manner
 - e.g., ‘buffer overrun’—a data structure **overflows its allocation**
 - **Time**—CPU interacts with resources no longer validly, e.g.,
 - ‘use after free’—a resource that **was deallocated is used again**
 - ‘time of check to time of use’ (TOCTOU)—the checking of a security condition **is decoupled from its use**
- Both approaches can affect **code** and/or **data**
 - CPU ends up running / reading / writing software it shouldn’t

Memory model for programming


- CPU's machine code **needn't be procedure-based** e.g.
 - can (potentially conditionally) jump to (goto) any other code
- Programming languages (PLs) usually more structured:
 - **Stack**: FIFO; memory lifecycle linked to PL functions' lifecycle
 - **Heap**: memory lifecycle decoupled from program flow
- CPUs very likely to support **call stack** explicitly
 - e.g., dedicated CPU registers for managing stack frames...

Heap-based (space) attacks on memory

- Consider a C program that `malloc`'s two 16 byte arrays
 - These arrays will be **allocated on the heap** (libc request of OS)
 - Let's pretend the arrays are allocated in sequential addresses
- Buffer overflow attack involves: (oversimplified)
 - If read/write index to first array **not bounds-checked** to be <16
 - ... then reads/writes on first array actually affect second array
 - (Real attacks often overwrite program code rather than data)
 - `strcpy` copies C strings without bounds check; use `strlcpy` !

Stack-based (space) attacks on memory

- Call stack pertains to code & data
 - Data: **local variables, parameters**
 - Code-relevant: **return address**
- Buffer overflow a local variable?
 - will potentially **rewrite return address**
 - function return **hands control to attacker's code**
- Note: stack address growth direction is CPU-dependent
 - x86 grows downwards (overflow of locals will reach return addr.)



Operating system (OS) memory protection

- Most CPUs have a **memory management unit (MMU)**
 - (Exception: old architectures and smaller embedded systems)
- MMU implements **usage restrictions on memory pages**
 - Pages are often 4KiB blocks of memory
 - Effects isolation of different operating system processes
 - Also separates applications from underlying OS kernel
- **Privilege escalation** attacks: into kernel from user code

Stop code being executed from data pages

- Execution space protection: **split code/data memory**
 - Prevents data pages having code executed from them
 - Represents OS+CPU increasingly **locking down use of memory**
 - Must couple with **address space layout randomisation** (ASLR)
- OS loader configures memory protection for app.
 - Application built with **clearly separated regions** (c.f., ELF)
- Mitigates attacks where **buffer overrun modifies code**
 - CPU capability: **NX** bit in AMD and **XD** bit in Intel CPUs

Code gadgets – attacks using existing code

- Execution space prot. stops attacker injecting code
 - ... however there's **already lots of code** on any target system
- Attacker can scan for '**gadgets**': abuse existing code
 - e.g., can jump into the middle of a destructive library function
 - may even be able to control parameters for those functions
- Significantly **raises the difficulty** of performing attacks
 - ... but many attackers are well resourced, and patient ...

Return oriented programming (ROP)

- Shown that **call stack attacks can chain gadgets**
- Attacker modifies return address and parameters
- Attacker isn't injecting code, just changing return addr.
- However net **effect is close to code injection**
- Not straightforward to detect attacks
 - solutions have proposed **integrity checks on return addresses**
 - ... but need to ensure that overheads are worth the expense

Vulnerabilities from parsing bugs

- **Parsing structured data** from simpler data can be risky:
 - Typical example, **SQL injection** (see next slide)
 - ... but also watch entities such as **file paths stored in strings**
- Many **forms of data parsing** are commonplace, e.g.:
 - **XML** documents
 - **Unicode** strings—e.g., UTF-8
 - **URLs**—e.g., wherein spaces are replaced by `%20` within in URLs
 - **Serialisation** of program objects (e.g., Java, Python, ...)

SQL injection attacks

- Code may build a string to submit query to database
- Form like `SELECT * FROM t WHERE t.name='$VAR'` is risky:
 - \$VAR needs to be checked to stop it escaping SQL statement
 - e.g., \$VAR should not contain single quotes
 - consider \$VAR being Robert'; DROP TABLE t (c.f. [XKCD comic 327](#))
- A solution: **SQL prepared statements** (? is placeholder)
 - `SELECT * FROM t WHERE t.name=?`—later bind variable to ?

File paths and potential security risks

- Common practice to store **file paths in string variables**
 - (Pain may be on offer regarding directory separator choice)
- Actually, **paths are far more subtle** than string suggests
 - Most operating systems can mount filesystems at any subpath
 - Independent parts of path string may be **case sensitive or not!**
- Another risk area: **simplification** using lexical processing
 - e.g., `thisDir/aDir/..../otherDir`—what if `aDir` is a symlink?

XML vulnerabilities

- Be careful parsing **untrusted source code**, here for XML:
 - **Entity attacks**, such as “Billion Laughs”:
 - `<!ENTITY lol9 "&lol8;&lol8;&lol8;&lol8;&lol8;&lol8;&lol8;"> etc.`
 - Very small file can easily explode to **occupy impractical resources**
 - External entity resolution: parser looks up **remote URLs / files**
 - **XSLT is Turing complete!** (XML stylesheet transformation)
- Use **existing XML parser** implementation, not your own
 - ... and continue installing its (likely many) **security updates**

Unicode handling can cause security issues

- Unicode: standard representing language characters
- UTF-8 is **variable-width** ASCII-compatible 8-bit encoding
 - Upper 128 values include mode shifts to multi-byte characters
 - Combining characters affect other characters:
 - e.g., <i U+00ED><diaeresis on previous U+0308> versus <i U+00EF>
 - normalisation is required to switch to longest / shortest form
- Security risk can relate to visual confusion or encoding
 - e.g., **normalisation failure** may lead to incorrect equality tests
 - Further risks **when embedded** in other forms, e.g., URI encoding

PL support for secure parsing

- Some PL syntax can **include XML directly**, e.g., Scala
 - Scala's parser accepts XML as the RHS of assignment:
 - `var myVariable = <p>A simple XML tree</p>`
 - (Scala is used in industry: e.g., LinkedIn, Twitter, Airbnb, Netflix, ...)
- PLs may also help security of **processing Unicode**:
 - e.g., Apple's Swift language ensures Unicode-correct handling
 - Would be non-idiomatic code to look at Swift string's bytes
- (Past project of mine added SQL&XML parsing to Python!)

Programming language choices for security

- Some situations require use of low-level, ‘unsafe’ PLs
 - e.g., directly **driving hardware devices** may need assembly
 - Most **mainstream OSs** have been largely coded in C/C++
- Applications can choose interpreted or compiled PLs
 - Security concerns are different, but **both have OS interfaces**
 - Compiled: result may have machine code vulnerabilities
 - Interpreted: likely rely on ‘foreign’ function interface (e.g., C)
 - e.g., Python often used as logic glue between C code libraries

Runtime support for managed PLs

- Manually performing **memory management** is riskier
 - ... although also necessary for corner cases
- Useful to seek **runtime systems that manage resources**
 - e.g., lifecycle of heap objects; OS interactions
 - Pay the ‘price’ of not directly controlling CPU with your code
- Java Virtual Machine: garbage collection; serialisation
 - JVM is now a target platform for other languages (e.g., Scala)

Functional programming languages

- Pure functional PLs **don't have intermediate state**
 - e.g., Haskell—variables are labels, not memory pigeonholes
 - ... but PLs have to **interact with underling OS** so pass state to fns:
 - FYI: ‘monads’ wrap functions and return values into efficient pipelines
- Many functional PLs are ‘impure’ with some state
 - i.e., state will likely involve mutable data structures
- DNS server ported to OCaml was more efficient than C:
 - ... it allowed safe reuse of memory where C code made copies

D (dlang) programming language

- Builds **pragmatic extensions** over C++
 - Bring desirable high-level functionality to low-level language
 - Aims to be as efficient as equivalent C++ but **terser and safer**
 - Still **supports inline assembly** (unlike C#, Java, etc.)
- Features that help security include integration of:
 - **Bounds-checked arrays**; garbage collection; strings are arrays
 - **@safe annotation** ensures valid lifetime of references
 - Compile-time check to **preclude use-after-free** types of errors
 - **'Better C' subset** removes D runtime, keeps bounds-checking, etc.

E—OO, secure, distributed PL

- Likely you'll never see/use E, but it is very well designed
- Method call = sending message to local/remote object
 - immediately—essentially like a function call (synchronous)
 - deferred—asynchronous, **caller gets ‘promise’** immediately
- **E objects are capabilities**, actually: controls visibility
 - Can use **sealer/unsealer pairs** to lock down object access
 - Can include guards to **check runtime conditions** ($\text{balance} \geq 0$)

Rust—low-level PL; more secure than C

- Began in Mozilla: e.g., for Servo secure browser (RIP)
 - Gaining adoption (but bumpy ride) in Linux kernel alongside C
- Key feature is the notion of **ownership typing**
 - If caller passes object to callee, caller can't modify it anymore
 - Rust 'borrow checker': ownership violations are **compiler errors**
 - (2024 Safe C++ proposal incorporates Rust borrow checker)
- Rust provides built-in **build system & package manager**

Engineering secure software

- Need security functionality? **Use existing libraries!**
 - e.g., NaCL; XACML; SAML; Kerberos GSS-API
 - ... you also need to **assess dependencies** and **apply updates**
- Apply defence in depth: **multiple layers of security**
 - Interacting with filesystem? Try to add a **chroot**
 - Handling sensitive data? Apply **encryption** defensively
 - Read-only database? Make **read-only replica**; don't trust code
 - Trade off (some) additional computing cost for extra security:
 - Use **short-lived OS processes** rather than risking memory leaks

In summary

- Described typical **machine model** and causes for **security problems** in time and space
- Outlined **machine code attack & defence evolution**
 - Discussed numerous common attack vectors
- Indicated how **choices of PL** can help security
 - New languages are still being developed...